

pipeline object
holds shaders and configuration
VkPipeline: Holds the state of the GPU needed to draw. For example: shaders, rasterization options, depth settings
A set of **VkPipeline** objects can be used for shader combinations and parameters needed to render materials

VkCommandPool
objects are used to allocate command buffers

render pass
VkRenderPass: Holds information about the images you are rendering into. All drawing commands have to be done inside a render pass. (main pass, shadow pass)

VkFramebuffer: Holds the target images for a render pass

A descriptor is a pointer to a resource (buffer or image). **VkDescriptorSet:** Holds the binding information that connects **shader inputs** to data such as **VkImage** textures and **VkBuffer** resources; a set of gpu-side pointers to be bound once. Descriptors must be grouped into sets, only sets can be bound.

VkInstance: The Vulkan context, used to access drivers

VkPhysicalDevice: A GPU. Used to query physical GPU details, like features, capabilities, memory size

VkDevice: The “logical” GPU context to execute commands on

command buffer
work GPU has to be recorded into a CommandBuffer, then submitted into a Queue
One command buffer / queue for each swap chain image
VkCommandBuffer: GPU commands
VkQueue: queue for commands

swap chain (list of images) display frame to the screen
VkSwapchainKHR: Holds the images for the screen. It allows you to render things into a visible window.
Has **VkQueue** handles

VkImage: A texture you can write to and read from (texture data)
VkImageView: representing contiguous ranges of the image sub-resources and containing additional metadata

VkBuffer: GPU visible memory (vertex data)

VkDescriptorSetLayouts: Blueprint for descriptor set, provided at pipeline creation
vkDescriptorPool: pre-allocated memory for descriptors, provided at pipeline creation

VkSemaphore: Synchronizes GPU to GPU execution of commands. Used for syncing multiple command buffer submissions one after another.

VkFence: Synchronizes GPU to CPU execution of commands. Used to know if a command buffer has finished being executed on the GPU

Render Loop

- request from the **VkSwapchainKHR** an image to render to (imageIndex)
- allocate a **VkCommandBuffer** from
 - a **VkCommandBufferPool**
 - reuse an already allocated command buffer that has finished execution
- write commands into command buffer
 - start a **VkRenderPass** (render into the image from swapchain, framebuffer[imageIndex])
 - for each object {
 - + bind a **VkPipeline**
 - + bind **VkDescriptorSet** resources for the shader parameters)
 - + bind the vertex buffers
 - + execute a draw call }
 - end the **VkRenderPass**
- end the **VkCommandBuffer**
- submit the command buffer into the queue for rendering
- present the image to the screen
- use a semaphore to make the presentation of the image wait until rendering is finished

Math Affine Transformation
Linear Transformations Scaling, Rotation around origin or anchor in 3x3 matrix, plus 4x1 translation (x,y,z,h coordinates)
h=1 for positions, h = 0 directions, distances, which are affected by scale&rotation, not translation
Euler angles will be used: Y(1), X(2), Z(3)
Quaternions have a few advantages over matrices (expressing 3D rotations)
Canonical view volume: normalized 2x2x1 **Vulkan:** y axis points down, z into screen (right hand) **Orthographic projection matrix:** scaled and translated from canonical volume, z useless except for depth test
Perspective projection: square frustum, orthographic matrix * perspective matrix
Viewport: specifies target framebuffer
Model matrix: TRS (from right to left: scale, rotate, translate) **Model-View-Projection matrix:** from left to right